

Convolution

Team 9

113062573 余伽璇

112078502 楊婷婷

Outline

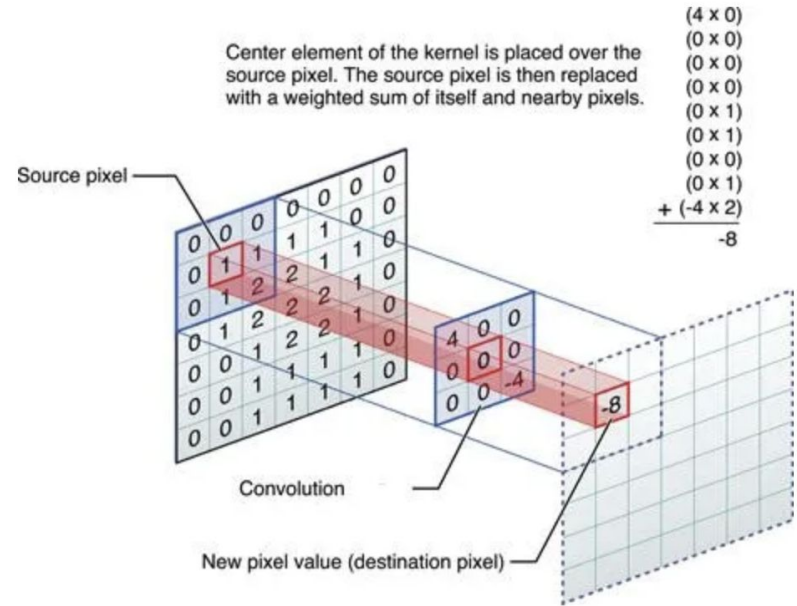
- Problem formulation
- Implementation
- Experiment
- Future work

Outline

- Problem formulation
- Implementation
- Experiment
- Future work

Convolution and Applications

Convolution is a fundamental operation in image processing, deep learning or signal processing etc. It involves the element-wise multiplication of two functions, one typically being a signal or image, and the other a kernel or filter. This operation captures local relationships and patterns, making it a powerful tool for feature extraction, system analysis.



Convolution Operation on a 7×7 matrix with a 3×3 kernel

Outline

- Problem formulation
- Implementation
 - sequential with vectorization
 - OpenMP
 - Pthread
 - MPI
 - Hybrid
 - single-GPU (regular version: global v.s shared)
- Experiment
- Future work

Implementation - Vectorization

- version 1
 - 256 bits-registers in AVX
 - $256 / 32 = 8$, 8 floats can be computed simultaneously
- version 2
 - Same as version 1, but using two `__m256` registers
 - 16 floats can be computed simultaneously

<i>Dist</i>			<i>Mask</i>		
D_1	D_2	D_3	M_1	M_2	M_3
D_4	D_5	D_6	M_4	M_5	M_6
D_7	D_8	D_9	M_7	M_8	M_9

Result

$$\begin{aligned} &+= D_1 \times M_1 + D_2 \times M_2 + D_3 \times M_3 + D_4 \times M_4 \\ &+ D_5 \times M_5 + D_6 \times M_6 + D_7 \times M_7 + D_8 \times M_8 \end{aligned}$$

Implementation - OpenMP

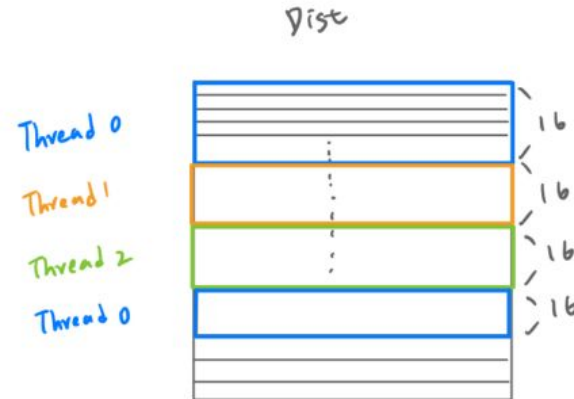
- Dynamic Scheduling

- `#pragma omp parallel for`
`schedule(dynamic, 16)`
- height partitioning



- Static Scheduling

- `#pragma omp parallel for`
`schedule(static, 16)`
- height partitioning



Implementation - Pthread

- Static Scheduling

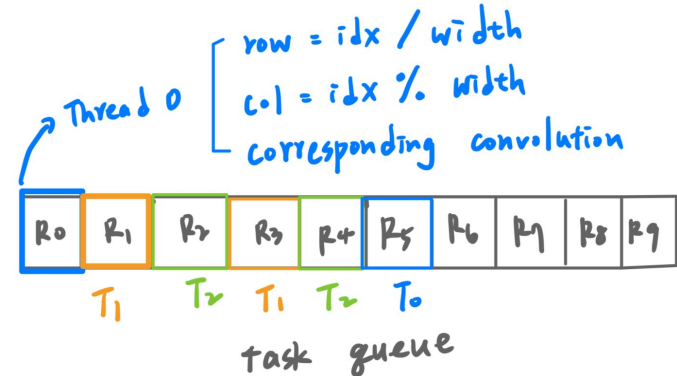
- divides the total number of elements in the output matrix equally among threads, with each thread handling a contiguous range of elements based on its thread ID

- Dynamic Scheduling

- master thread
- task_queue : store #(result matrix size) tasks and one termination signal (-1) at the end
- mutex lock

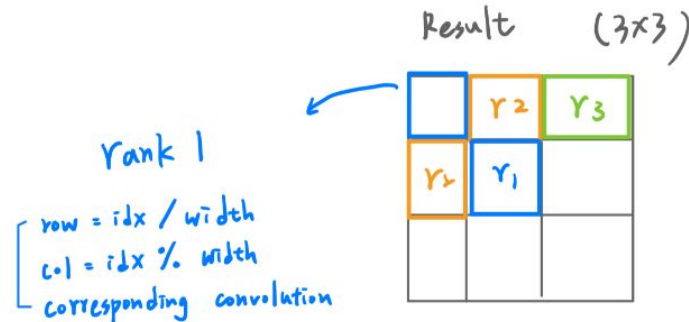
Result (3x3), Thread x4

T ₀	T ₀	T ₁
T ₁	T ₂	T ₂
T ₃	T ₃	T ₃



Implementation - MPI

- master process : rank 0
- send -1 to work process if all tasks are done
 - Ex : 4 process, input = 5*5, kernel = 3*3, result = 3*3



- How the master process collects result data ?

```
MPI_Recv(&task_index, 1, MPI_INT, MPI_ANY_SOURCE, 1, MPI_COMM_WORLD, &status);  
MPI_Recv(&result, 1, MPI_FLOAT, status.MPI_SOURCE, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
  
Result[task_index] = result;
```

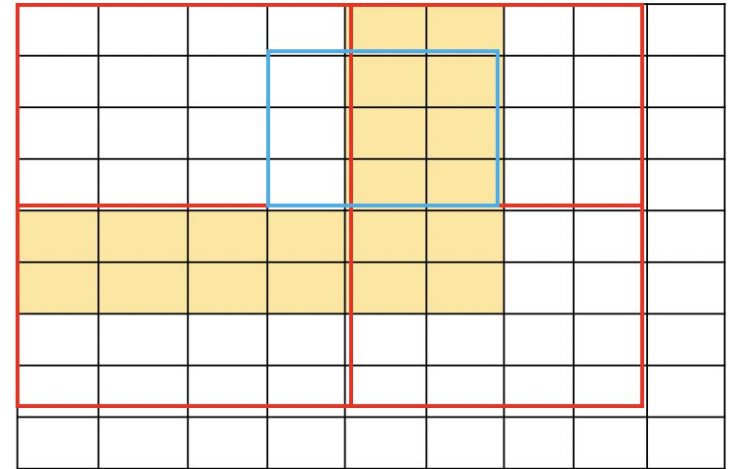
Implementation - Hybrid

- Task partition and process communication same as MPI
- Convolution of each work process use :

```
#pragma omp simd reduction(+:sum)
for (ki = 0; ki < k; ++ki)
{
    for (kj = 0; kj < k; ++kj)
    {
        sum += Dist[(row + ki) * width + (col + kj)] * Mask[ki * k + kj];
    }
}
```

CUDA - Regular Convolution

- Global Memory:
all compute fetches from global, no boundary issue
- Share Memory :
 - need **apron** to load additional data to the thread block
 - e.g when threadblock = 4×4 , kernel = 3×3 at block (0,0) need the data from the other threadblocks→ need additional data in share mem



Outline

- Problem formulation
- Implementation
- Experiment
- Future work

Experiment - Set up

- Environment of CPU

```
pp24s111@nthu-master:~/final$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Address sizes:          52 bits physical, 57 bits virtual
Byte Order:             Little Endian
CPU(s):                 192
On-line CPU(s) list:    0-191
Vendor ID:              GenuineIntel
Model name:             INTEL(R) XEON(R) PLATINUM 8568Y+
CPU family:             6
Model:                  207
Thread(s) per core:     2
Core(s) per socket:     48
Socket(s):              2
Stepping:               2
BogoMIPS:               4600.00
Flags:                  fpu vme de pse tsc msr pae mce cx8 a
nonstop_tsc cpuid aperfmperf tsc_known
d lahf_lm abm 3dnowprefetch cpuid_fault
rms invpcid cqm rdt_a avx512f avx512
ni avx512_bf16 wbnoinvd dtherm ida a
ovdirt64b enqcmd fsrm md_clear serial
Virtualization features:
  Virtualization:       VT-x
Caches (sum of all):
  L1d:                   4.5 MiB (96 instances)
  L1i:                   3 MiB (96 instances)
  L2:                   192 MiB (96 instances)
  L3:                   600 MiB (2 instances)
NUMA:
  NUMA node(s):          2
  NUMA node0 CPU(s):     0-47,96-143
  NUMA node1 CPU(s):     48-95,144-191
Vulnerabilities:
  Gather data sampling:   Not affected
```

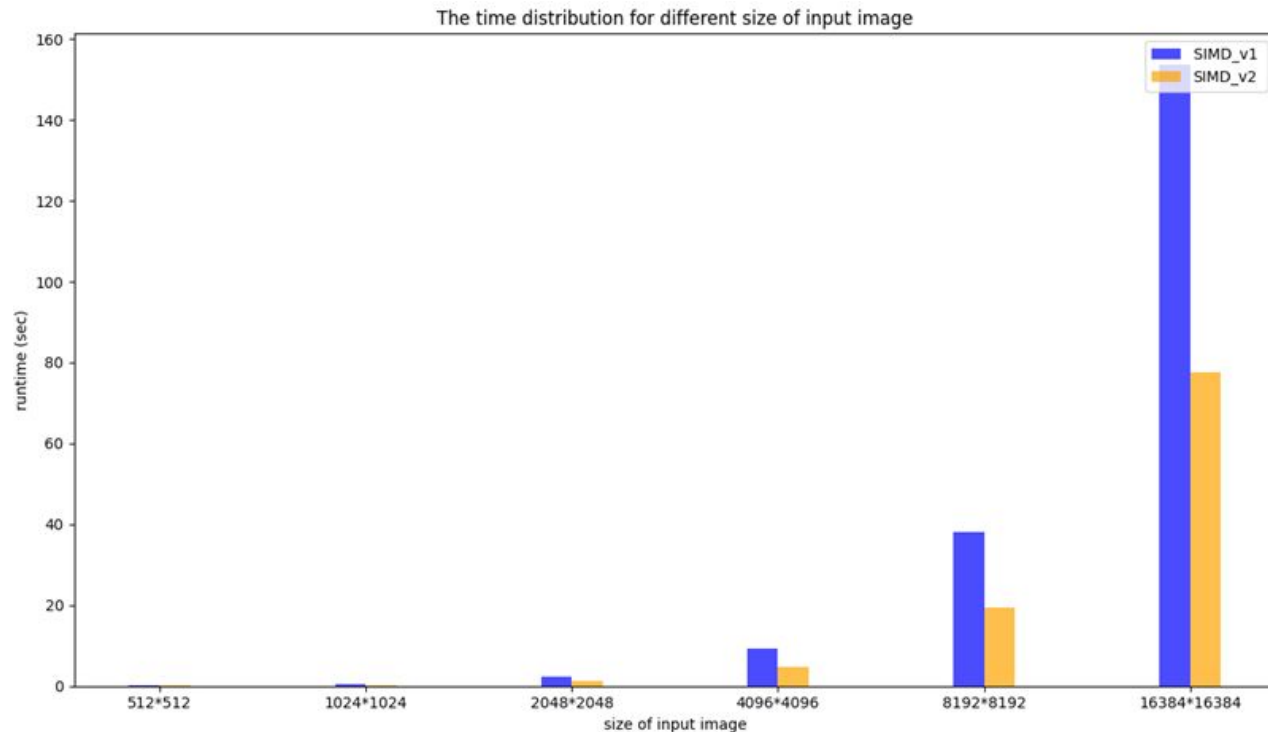
- Environment of GPU

```
Device 0: "NVIDIA GeForce GTX 1080"
  CUDA Driver Version / Runtime Version      12.6 / 12.6
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             8107 MBytes (8500871168 bytes)
  (20) Multiprocessors, (128) CUDA Cores/MP: 2560 CUDA Cores
  GPU Max Clock rate:                        1835 MHz (1.84 GHz)
  Memory Clock rate:                         5005 Mhz
  Memory Bus Width:                          256-bit
  L2 Cache Size:                             2097152 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:           65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                 32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:       Yes with 2 copy engine(s)
  Run time limit on kernels:                  No
  Integrated GPU sharing Host Memory:         No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                     Disabled
  Device supports Unified Addressing (UVA):    Yes
  Supports Cooperative Kernel Launch:         Yes
  Supports MultiDevice Co-op Kernel Launch:   Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 1
  Compute Mode:
    |< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.6, CUDA Runtime Version = 12.6, NumDevs = 1
Result = PASS
```

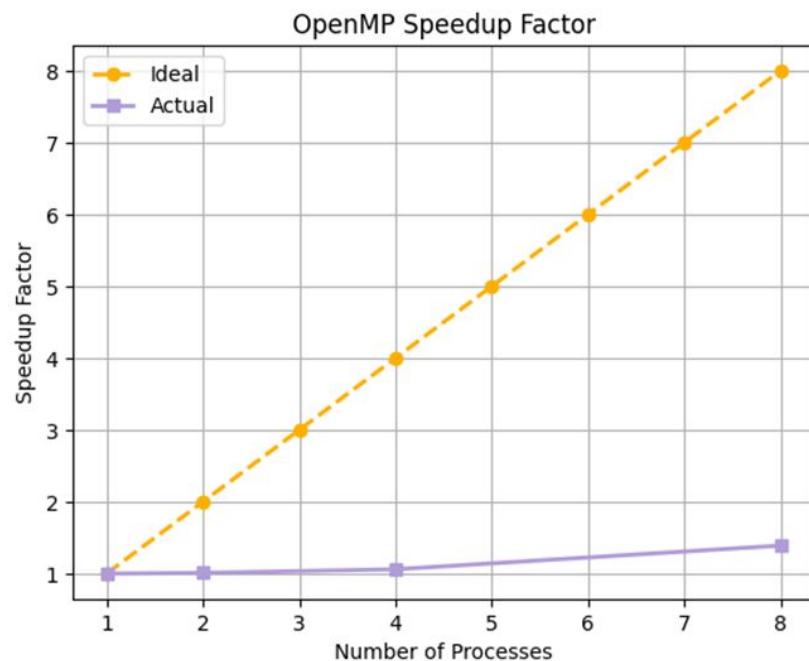
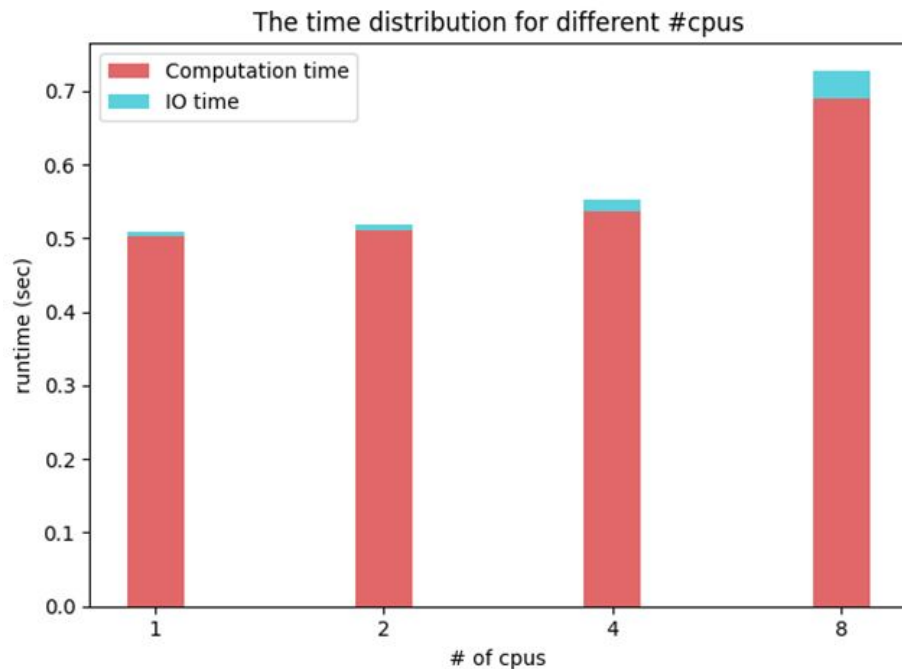
Experiment - Vectorization

- Sequential version with SIMD-v1 vs SIMD-v2
 - kernel size : 96*96



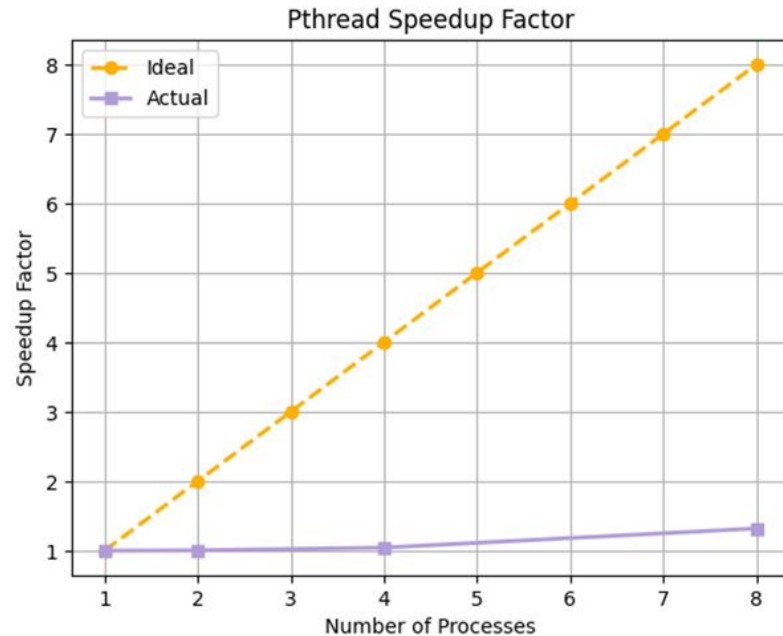
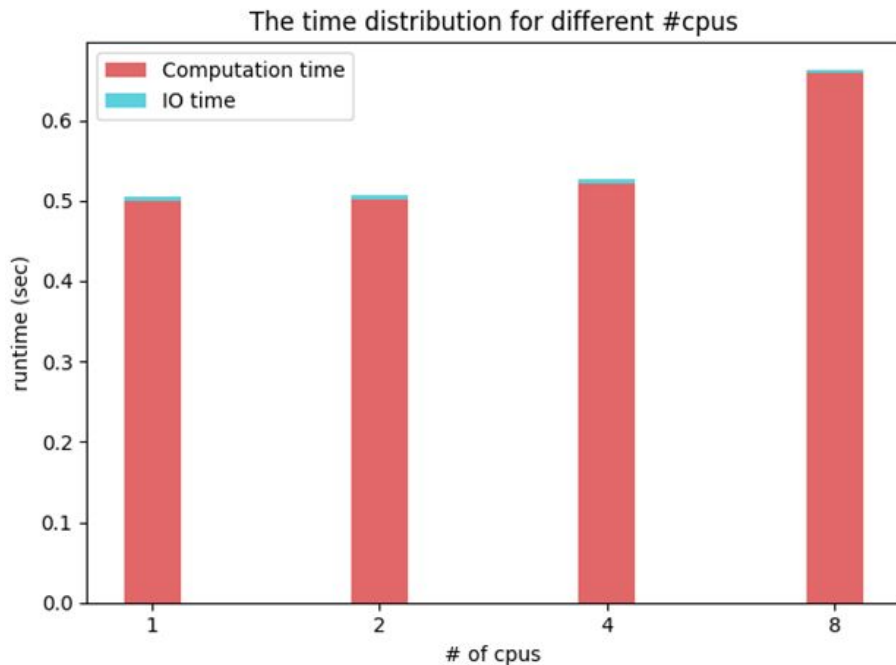
Experiment - OpenMP

- static scheduling with SIMD-v1
 - testcase : input size = 1024×1024 , kernel size = 96×96



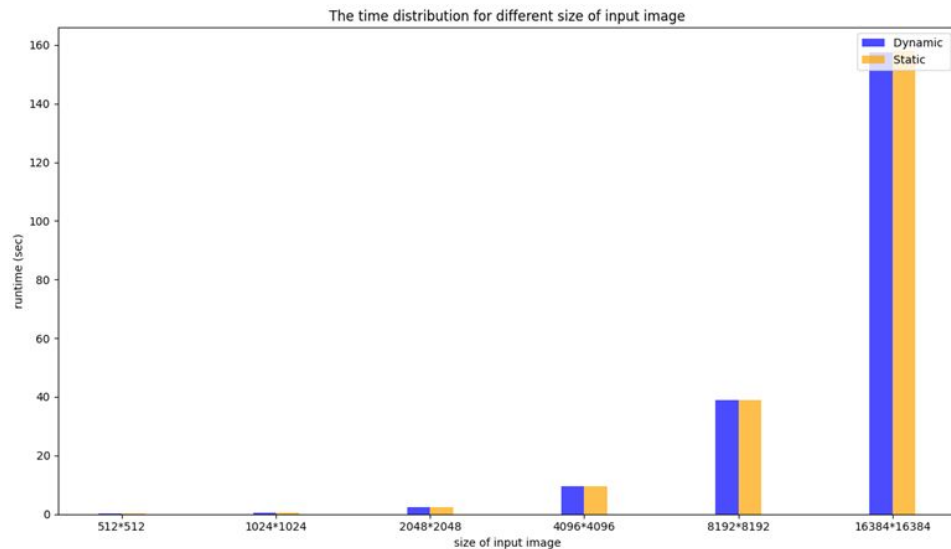
Experiment - Pthread

- static scheduling with SIMD-v1
 - testcase : input size = 1024×1024 , kernel size = 96×96

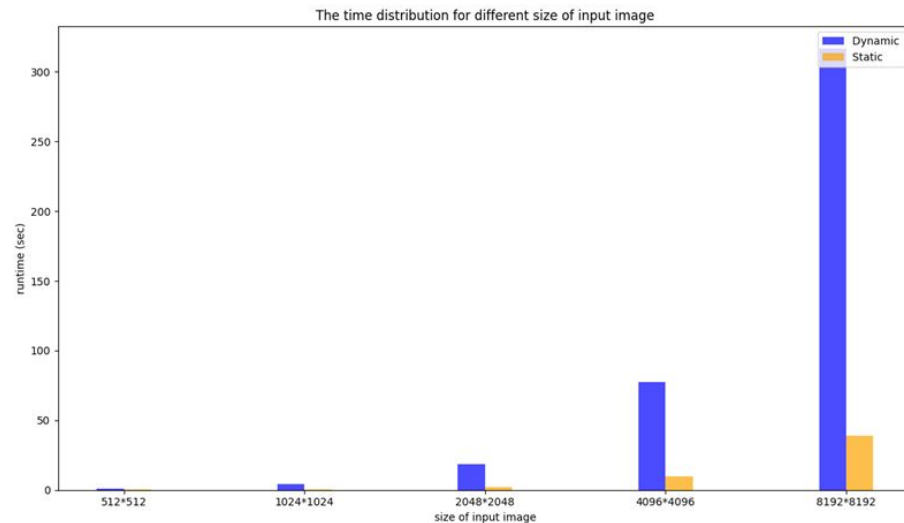


Experiment - OpenMP & Pthread

- Scheduling method : dynamic , static
 - kernel size = 96*96, 4 cores



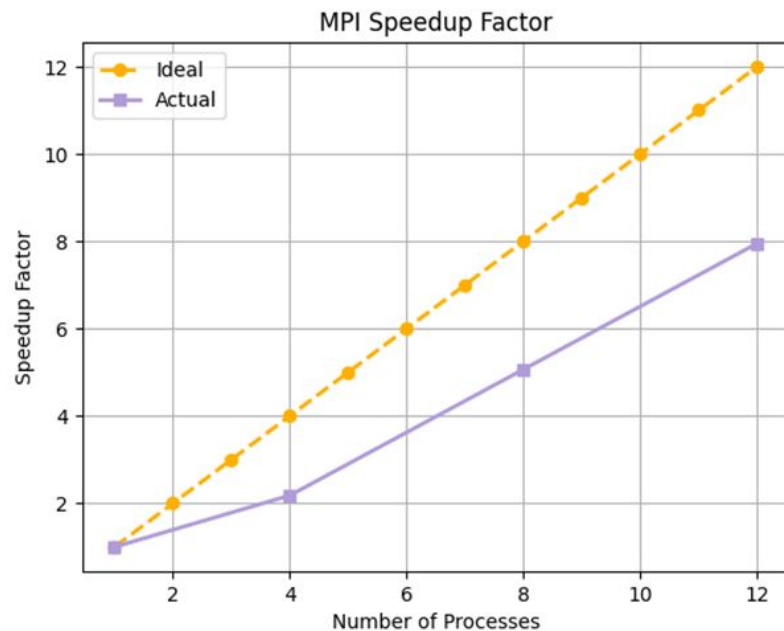
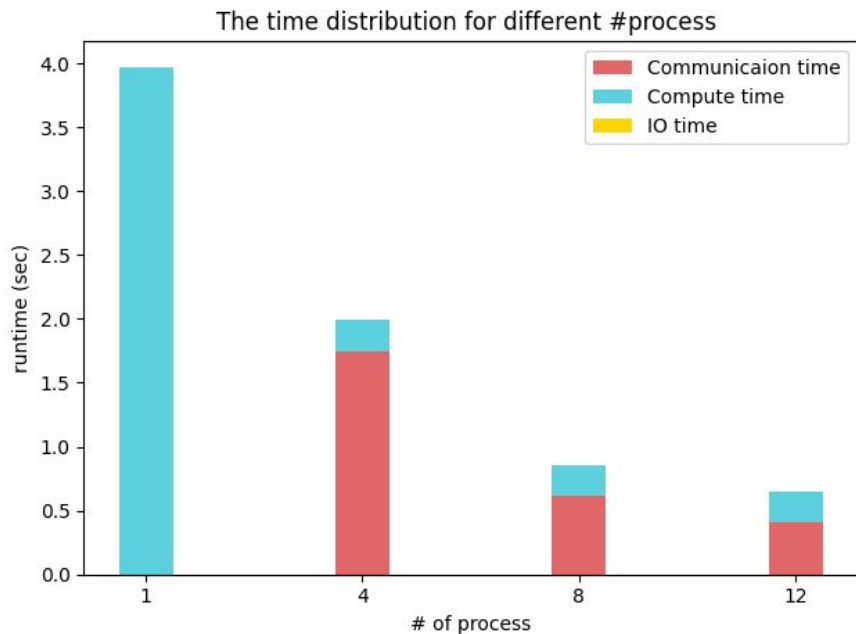
OpenMP



Pthread

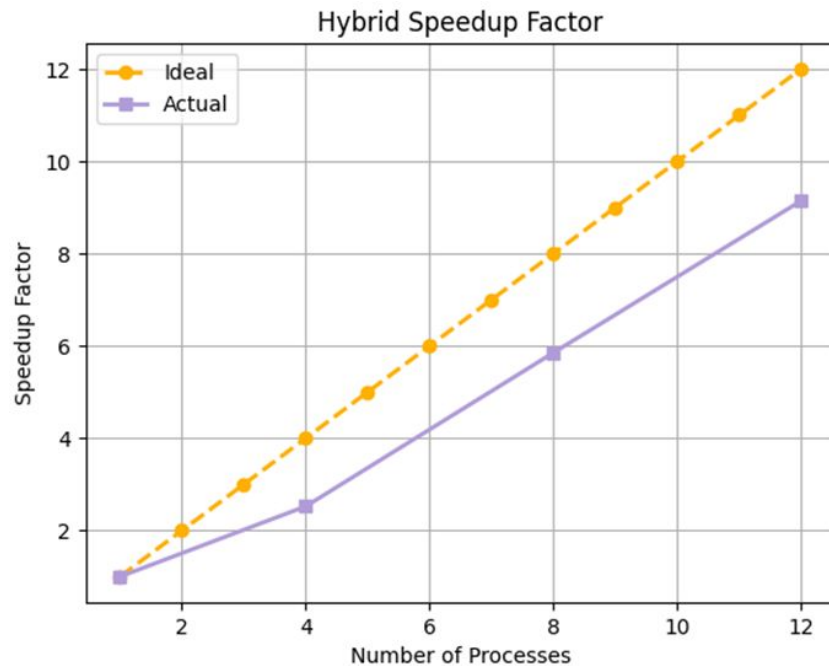
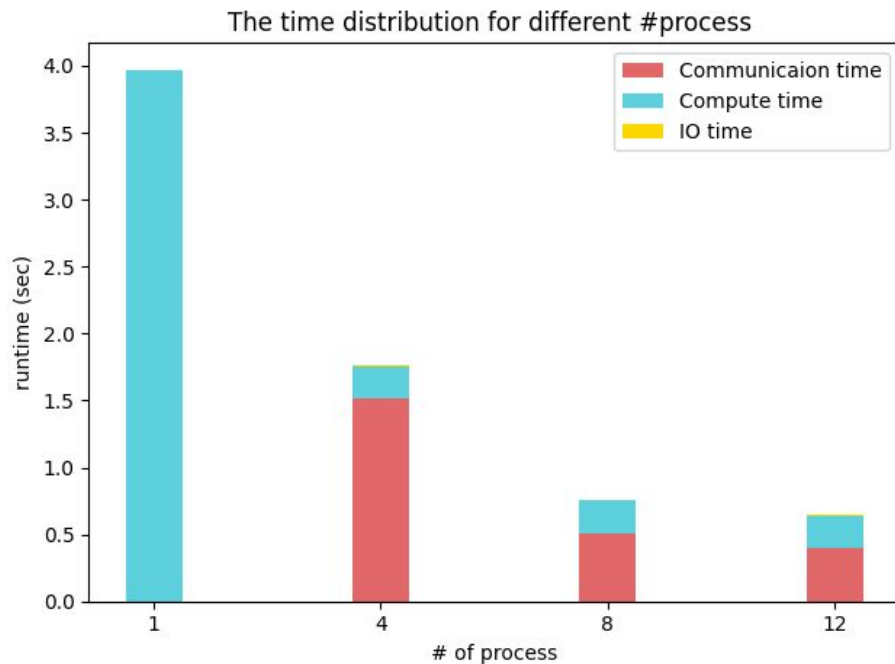
Experiment - MPI

- 4 nodes
 - testcase : input size = 1024×1024 , kernel size = 96×96



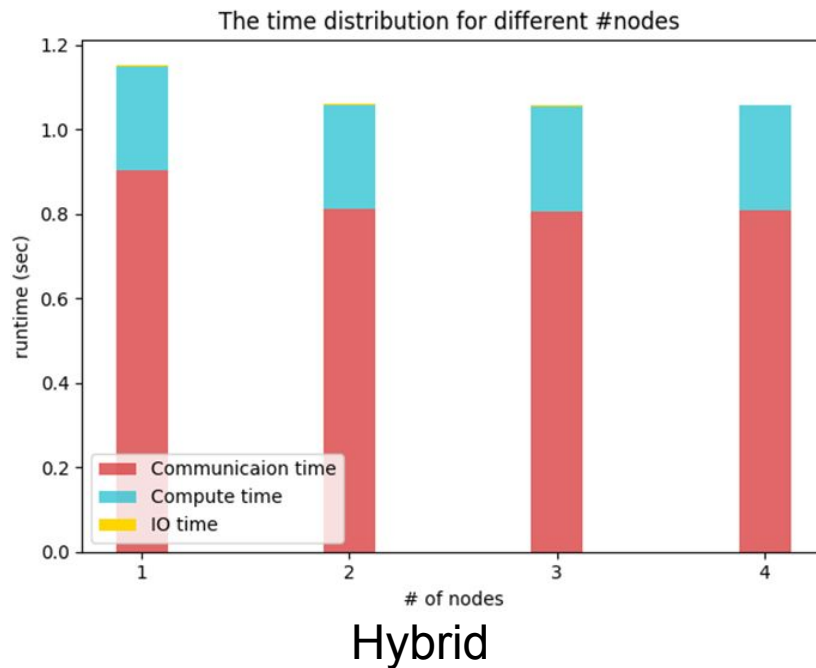
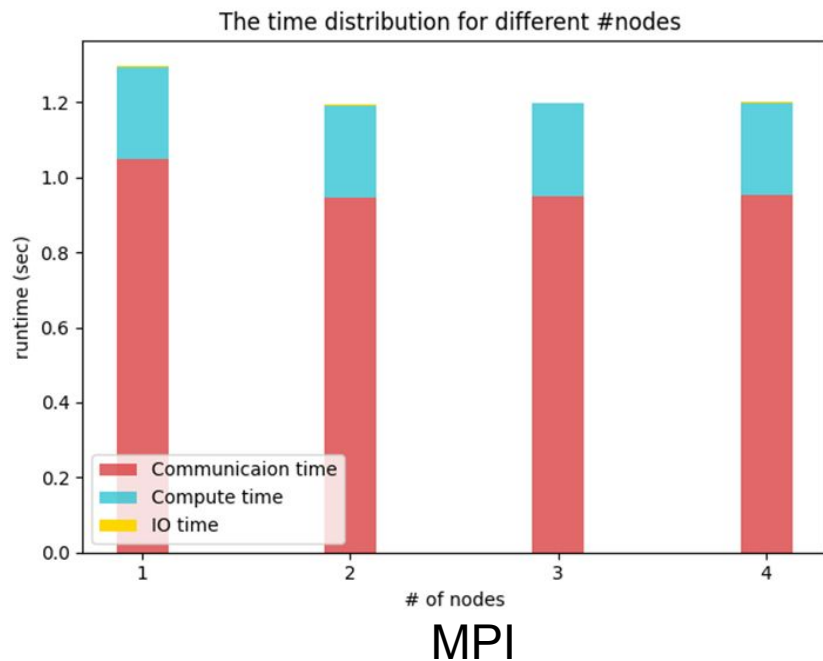
Experiment - Hybrid

- 4 nodes
 - testcase : input size = 1024×1024 , kernel size = 96×96



Experiment - MPI & Hybrid

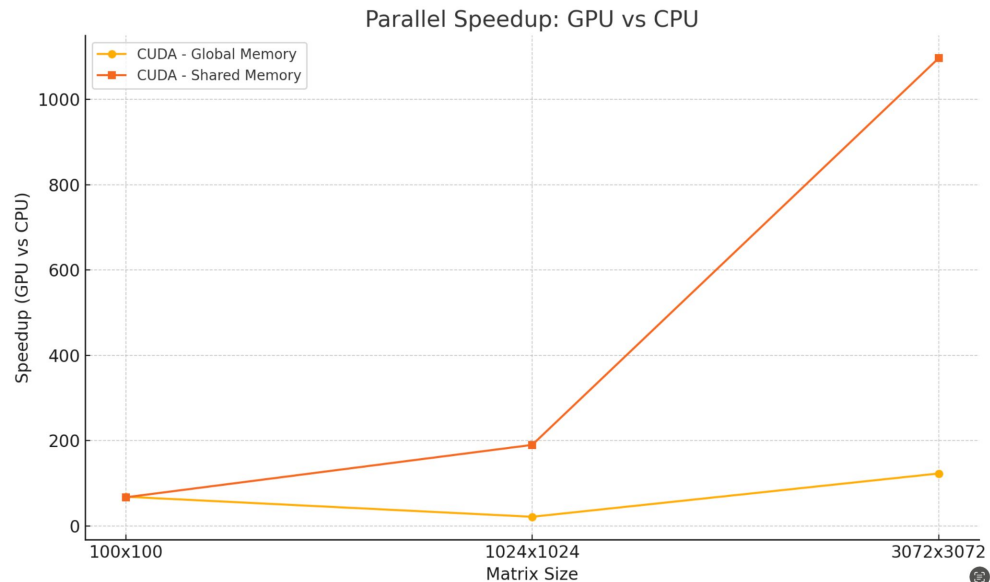
- Comparison of different #nodes
 - testcase : input size = 1024×1024 , kernel size = 96×96 , ppn = 6



Experiment - Speedup GPU v.s sequential CPU

Compute Time (ms)	100 100 5	1024 1024 96	3072 3072 17
CPU - naïve sequential	0.200	3,932.212	3125.228
CUDA - global memory	0.104	56.326	39.733
CUDA - shared memory	0.107	0.110	4.4482
CPU - parallel version (MPI)	7.124	1197	4878.624

Speedup / against GPU	100 100 5	1024 1024 96	3072 3072 17
CUDA - global memory	1.915708812	69.812	78.656
CUDA - shared memory	1.877934272	35877.847	702.58
Speedup / against GPU -global	100 100 5	1024 1024 96	3072 3072 17
CUDA -speed up smem / global	0.98028169	513.923	8.932
Speedup / against CPU - parallel	100 100 5	1024 1024 96	3072 3072 17
CUDA - global memory	68.2375	21.2513	122.7852
CUDA - shared memory	66.8920	189.8240	1096.7636



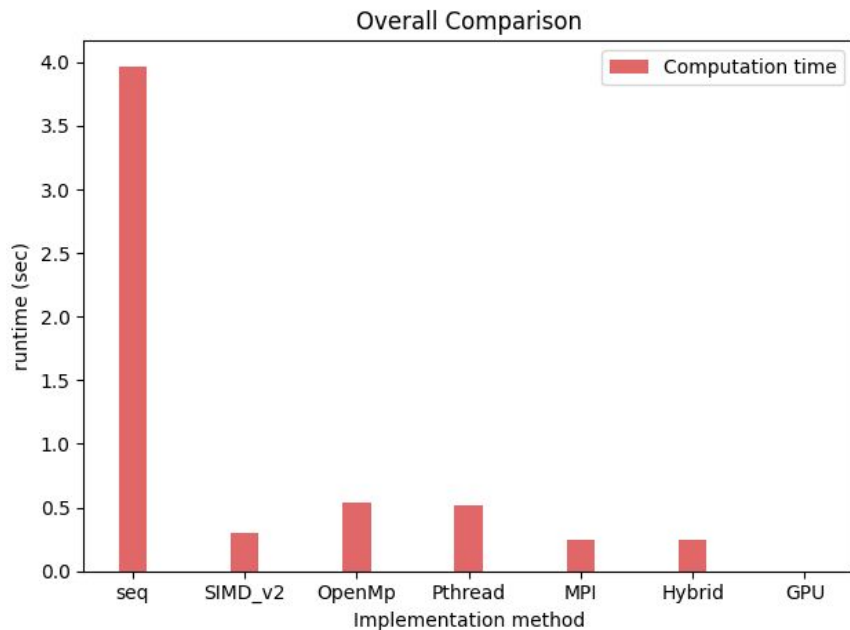
GPU Analysis

```
==1714625== Profiling application: ./conv_gpu_v2 ./testcases/1024_1024_96.in ./profile.out
==1714625== Profiling result:
   Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities: 56.75% 802.19us      2 401.09us 4.3520us 797.83us [CUDA memcpy HtoD]
                43.25% 611.40us      1 611.40us 611.40us 611.40us [CUDA memcpy DtoH]
API calls:      97.32% 104.02ms      3 34.673ms 64.887us 103.88ms cudaMalloc
                1.89% 2.0173ms      3 672.44us 79.650us 1.0278ms cudaMemcpy
                0.40% 425.97us      3 141.99us 93.392us 187.82us cudaFree
                0.17% 180.51us     114 1.5830us 92ns 70.661us cuDeviceGetAttribute
                0.11% 120.28us      1 120.28us 120.28us 120.28us cudaLaunchKernel
                0.04% 40.553us      2 20.276us 620ns 39.933us cudaEventCreate
                0.03% 27.998us      2 13.999us 3.7630us 24.235us cudaEventRecord
                0.02% 17.822us      1 17.822us 17.822us 17.822us cuDeviceGetName
                0.02% 16.323us      1 16.323us 16.323us 16.323us cuDeviceGetPCIBusId
                0.01% 12.426us      1 12.426us 12.426us 12.426us cudaEventSynchronize
                0.00% 4.3800us      1 4.3800us 4.3800us 4.3800us cudaEventElapsedTime
                0.00% 1.1980us      3 399ns 131ns 867ns cuDeviceGetCount
                0.00% 554ns      2 277ns 119ns 435ns cuDeviceGet
                0.00% 356ns      1 356ns 356ns 356ns cuDeviceTotalMem
                0.00% 342ns      1 342ns 342ns 342ns cuModuleGetLoadingMode
                0.00% 226ns      1 226ns 226ns 226ns cuDeviceGetUuid
```

- CGMA ratio
- align shared memory

Experiment - Overall Comparison

- testcases : input size = 1024×1024 , kernel size = 96×96
- Only consider the computing time and communication time
 - Sequential
 - SIMD_v2
 - OpenMP + SIMD_v1 + 4 cores
 - Pthread + SIMD_v1 + 4 cores
 - MPI + 4 nodes + (ppn=12)
 - Hybrid + 4 nodes + (ppn=12)
 - GPU (shared memory)



Outline

- Problem formulation
- Implementation
- Experiment
- Future work

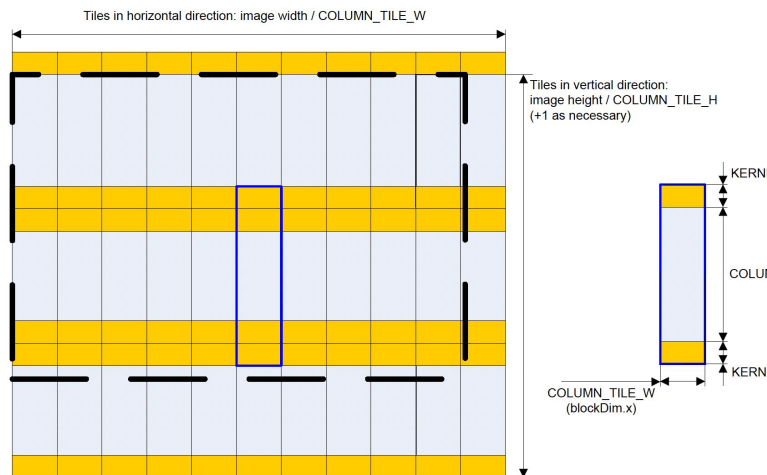
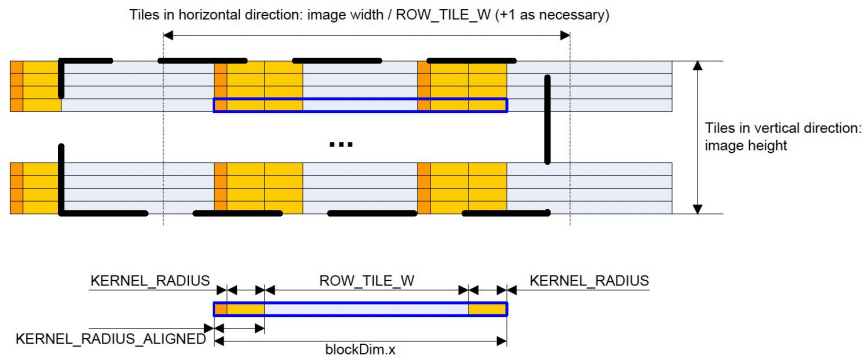
Future work: Separatable Convolution

- Devide into two one dimensional filter (row convolution -> col convolution)
- Constraint: kernel matrix must be rank-1
→ so can be written as two rank-1 kernel matrix outer product

Applying $\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$ to the data is the same as applying $\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$ followed by $\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$.

- Application example of kernel matrix are rank-1 : Gaussian Blur, Sobel Edge detect (lab practice)

How it works and Application



- Good for big kernel size
- Time complexity:

$$O(n^2 \cdot k^2) \longrightarrow O(n^2 \cdot k)$$

- save share memory usage, store k , instead of $k \cdot k$ in each pass (can process the real data)
- Better scalability
- direct: growth quadratic k^2
- seperatable: growth lineary k

Thanks