

CS542200 Parallel Programming

Homework 5: Observe the behavior of UCX yourself

Due: Sun. Dec 25 23:59, 2024

1 GOAL

Develop observability code within the UCX framework to gain a comprehensive understanding of its architecture and enhance our knowledge in modern parallel computing frameworks, with the following objectives:

1. **Comprehend the UCX Framework:** Delve into the UCX framework to understand the roles and interconnections of its subcomponents, thereby gaining a deeper insight into its overall functionality.
2. **Examine TLS Impact:** Investigate how Transport Layer Security (TLS) is influenced by the interactions and dependencies between UCP, UCS, and UCT within the UCX framework.
3. **Learn Program Structure:** Acquire knowledge about the program structure of modern parallel computing frameworks, focusing on how they are designed, implemented, and optimized for performance.

2 REQUIREMENTS

-
- In this assignment, you are required to modify and recompile the UCX system (the underlying communication protocol for MPI) to add the following two features:
 - Display the UCX transport protocols currently configured in the system.
 - ◆ You are required to insert code at `// TODO: PP-HW (src/ucs/config/parser.c)` to ensure your program is compatible with the UCX framework.
 - Output the UCX transport protocols actively utilized by the program.

3 RUN YOUR PROGRESS

-
1. Recompile UCX and install it in the `$HOME/ucx-pp` directory, ensuring that a `lib` folder is present that contains the required dynamic link files. (`libucp.so.0`, `libuct.so.0`, `libucm.so.0`, `libucs.so.0`).

Note that every time you modify the UCX code, you need to re-execute the last line of commands to recompile.

```
git clone -b pp2024 https://github.com/NTHU-LSALAB/UCX-lsalab
cd UCX-lsalab
mkdir build && cd build
../configure --prefix=$HOME/ucx-pp/ --with-go=no
srunk -n 1 -c 12 make -j12 install
```

2. We provide a custom command `mpiucx`, designed to replace `mpirun`, which will utilize your self-compiled UCX as the underlying communication framework for running MPI parallel programs. You can find the test cases in `UCX-lsalab/test/`

```
module load openmpi/ucx-pp
mpiucx -x UCX_LOG_LEVEL=info -np 2 ./mpi_hello.out
```

3. Alter the internal code of UCX to enable printing of:
 - The UCX_TLS information is currently specified by the system.
 - The final TLS transport method was selected by UCX.
4. You may need to trace the files within `src/ucp`, `src/uct`, and `src/ucs` directories and modify at least the following two files to complete the homework.
 - `ucp/core/ucp_worker.c`
 - Invoke `ucp_config_print` to print UCX_TLS.
 - Print Line 2 by identifying a key variable.
 - `ucs/config/parser.c`
 - Implement the `ucs_config_parser_print_opts` function to be invoked by `ucp_config_print` here.

4 RESULT

For each transport, the output comprises 2 lines.

- Line 1: The information from UCX_TLS needs to be exactly the same.
- Line 2: The information should include only the key strings of transport protocols selected by UCX. The example must contain strings.

```
cfg#0 tag(sysv/memory cma/memory)
```

Sample input & output

- For the `openmpi/ucx-pp` configuration, the default transport protocol is set to utilize `ud_verbs`.

```
[willian@apollo31 mpi]$ mpiucx -np 1 ./mpi_hello.out
UCX_TLS=ud_verbs
0x56544467c8f0 self cfg#0 tag(ud_verbs/ibp3s0:1)
Hello world from processor apollo31, rank 0 out of 1 processors
```

- We have enabled UCX to leverage all available transport protocols by setting up `UCX_TLS`.

```
[willian@apollo31 mpi]$ mpiucx -n 2 -x UCX_TLS=all ./send_recv.out
UCX_TLS=all
0x5580604e1a90 self cfg#0 tag(self/memory cma/memory)
UCX_TLS=all
0x559a90bb7a00 self cfg#0 tag(self/memory cma/memory)
UCX_TLS=all
0x5580604e1a90 intra-node cfg#1 tag(sysv/memory cma/memory)
UCX_TLS=all
0x559a90bb7a00 intra-node cfg#1 tag(sysv/memory cma/memory)
Process 0 sent message 'Hello from rank 0' to process 1
Process 1 received message 'Hello from rank 0' from process 0
```

5 REPORT

Here is the report template <https://hackmd.io/@William-Mou/pp-hw5-report>

Please download the markdown and submit the .md file.

6 GRADING

1. [30%] Correctness

- The pass rate of test data for ucx-judge.

2. [20%] Demo

- A demo session will be held remotely. You'll be asked questions about the homework.

3. [50%] Report

- Grading is based on your evaluation, discussion and writing. If you want to get more points, design or conduct more experiments to analyze your implementation.

7 SUBMISSION

Upload the files below to eeclass. (**DO NOT COMPRESS THEM**)

- hw5.diff
 - You can obtain this file by
 - `git add -A && git commit -m "Observed the behavior of UCX."`
 - `git diff --color remotes/origin/pp2024 > hw5.diff`
- hw5_{student_ID}.md or hw5_{student_ID}.pdf
 - If possible, please submit it in Markdown and provide a HackMD link or PDF file as a backup.
 - Please confirm that your image is displayed correctly.
- **KEEP \$HOME/ucx-pp/ CORRECTLY**

8 HINTS

- Type `ucx-judge` on **Apollo** to run the test cases.
- Scoreboard:
 - <https://apollo.cs.nthu.edu.tw/pp24/scoreboard/ucx/>
- You can use:
 - `ucx_info -d` to show all available TLS in your cluster.
 - `-x UCX_LOG_LEVEL=info` in `mpirun` to obtain debug information during UCX runtime
 - `-x UCX_TLS=<TLS>` to specify the permitted TLS for your program.
- Refer to **slide p37** of the presentation to determine which component is accountable for the UCX_TLS configuration in the system.
- Consult **slides p41 to p42** of the presentation for insights into UCP's architecture, which will aid in navigating your code to pinpoint the transport protocols actually in use.
- If you encounter any problems with the homework specification, judge scripts, example source code, or test cases, please contact the TA at pp@lsalab.cs.nthu.edu.tw or eeclass.

- You are allowed to discuss and exchange ideas with others, but you are required to write the code yourself. If we find you cheating, you'll receive 0 points.